

# 50.042 FCS Summer 2024

## Lecture 5 – Applications for Hashing

Felix LOH

Singapore University of Technology and Design



SINGAPORE UNIVERSITY OF  
TECHNOLOGY AND DESIGN

With selected materials adapted from: *Understanding Cryptography: A Textbook for Students and Practitioners*, by C. Paar and J. Pelzl

# Applications for cryptographic hash functions

- Hash functions implement a one-way, non-invertible, function on data  
*→ Mathematically impossible to find inverse.*
  - Contrast these with encryption/decryption functions, which must be invertible (or bijective) functions
  - Ideally, the hash output (i.e. hash value) does not reveal any information about the input
- Hash functions can be used for:
  - Commitment schemes
  - Protection of message integrity and authentication
  - Error detection
  - Storage of secrets (with some caveats)

# A commitment scheme

- Motivating problem: Suppose we have several individuals bidding on an item in a sealed bid scheme:
  - Each person can submit one sealed bid
  - Bids are compared after everybody has submitted their bids
  - The person with the highest bid gets the item
- How can we collect the bids and determine the highest bid securely (without using shared keys)?
- Alternative problem: Alice and Bob play a game of rock, paper, scissors over a remote connection
  - Need to seal each player's choice securely, then only reveal the two choices simultaneously after both players have submitted their choice

# A commitment scheme

- Both problems can be solved using a cryptographic hash scheme  $z = h(x)$
- Use a two-phase protocol:
  1. Alice and Bob choose their respective actions e.g.  $x_a = \text{"scissors"}$  and  $x_b = \text{"rock"}$ , then they compute the corresponding hashes  $z_a = h(x_a)$  and  $z_b = h(x_b)$  and exchange their commitments  $z_a$  and  $z_b$
  2. Alice and Bob then exchange their actual messages  $x_a$  and  $x_b$ , and they each verify that the message they received is authentic by recomputing the hash  $h(x_a)$  or  $h(x_b)$  of the received message and comparing it with the corresponding commitment hash  $z_a$  or  $z_b$  that they received in the 1<sup>st</sup> phase

# An issue with the commitment scheme

- There is a problem with the commitment scheme as described in the previous slide:
  - During the 1<sup>st</sup> phase of the scheme (before the phase is completed), let's assume that Alice decides on some action and sends her commitment hash  $z_a$  over to Bob
  - Bob can then reverse-engineer Alice's commitment hash to identify the action that Alice chose
  - He then chooses the optimal action, then sends Alice his commitment hash  $z_b$

# An issue with the commitment scheme

- This reverse-engineering is possible because 'rock', 'paper' and 'scissors' always hash to the same values in the hash function
- There is only a **small set of possible actions** (i.e. three values), so it's easy for Bob to just precompute all three possible hash values to determine Alice's chosen action **before** she sends her message in the 2<sup>nd</sup> phase
- We need to make the **input less predictable** (i.e. increase the input space)

# Fixing the issue with the commitment scheme

- We can use *salting* to increase the input space of the hash function
- This basically entails adding some additional bits of data (i.e. a salt) to the original message *before* running the hash function
- The salt value can be different, even if the original message is the same



# Fixing the issue with the commitment scheme

- Modify the two-phase protocol as follows:
  1. Alice and Bob choose their respective actions along with some random salt value  $s$ , then they compute the corresponding hashes  $z_a = h(x_a, s_a)$  and  $z_b = h(x_b, s_b)$  and exchange their commitment hash  $z_a$  and  $z_b$
  2. Alice and Bob then exchange their message-salt pairs  $x_a, s_a$  and  $x_b, s_b$  and they each verify that the message they received is authentic by recomputing their respective hash  $h(x_a, s_a)$  or  $h(x_b, s_b)$  and comparing it with the respective commitment hash  $z_a$  or  $z_b$  that they received in the 1<sup>st</sup> phase

# Padding for hash functions

- Recall that from the last lecture that we carry out padding for the last block of the message, before passing the message blocks to the hash function
- E.g. for SHA-1, the padding method is a '1' followed by a bunch of zeroes, followed by the 64-bit binary representation of the length of the message in bits before padding
- Other hash functions and ciphers use different padding methods
- Note that we **cannot** just simply pad the last block with random bits, because then it is no longer possible to verify the hash output (it is no longer deterministic)

# Message authentication codes (MACs)

- Recall from last week's lectures that encryption (or ciphers) **cannot** guarantee message integrity
  - Recall the “buy100” example
  - Eve can still modify the message despite not knowing the secret key
- Hash functions can guarantee integrity to a certain extent

# Message authentication codes (MACs)

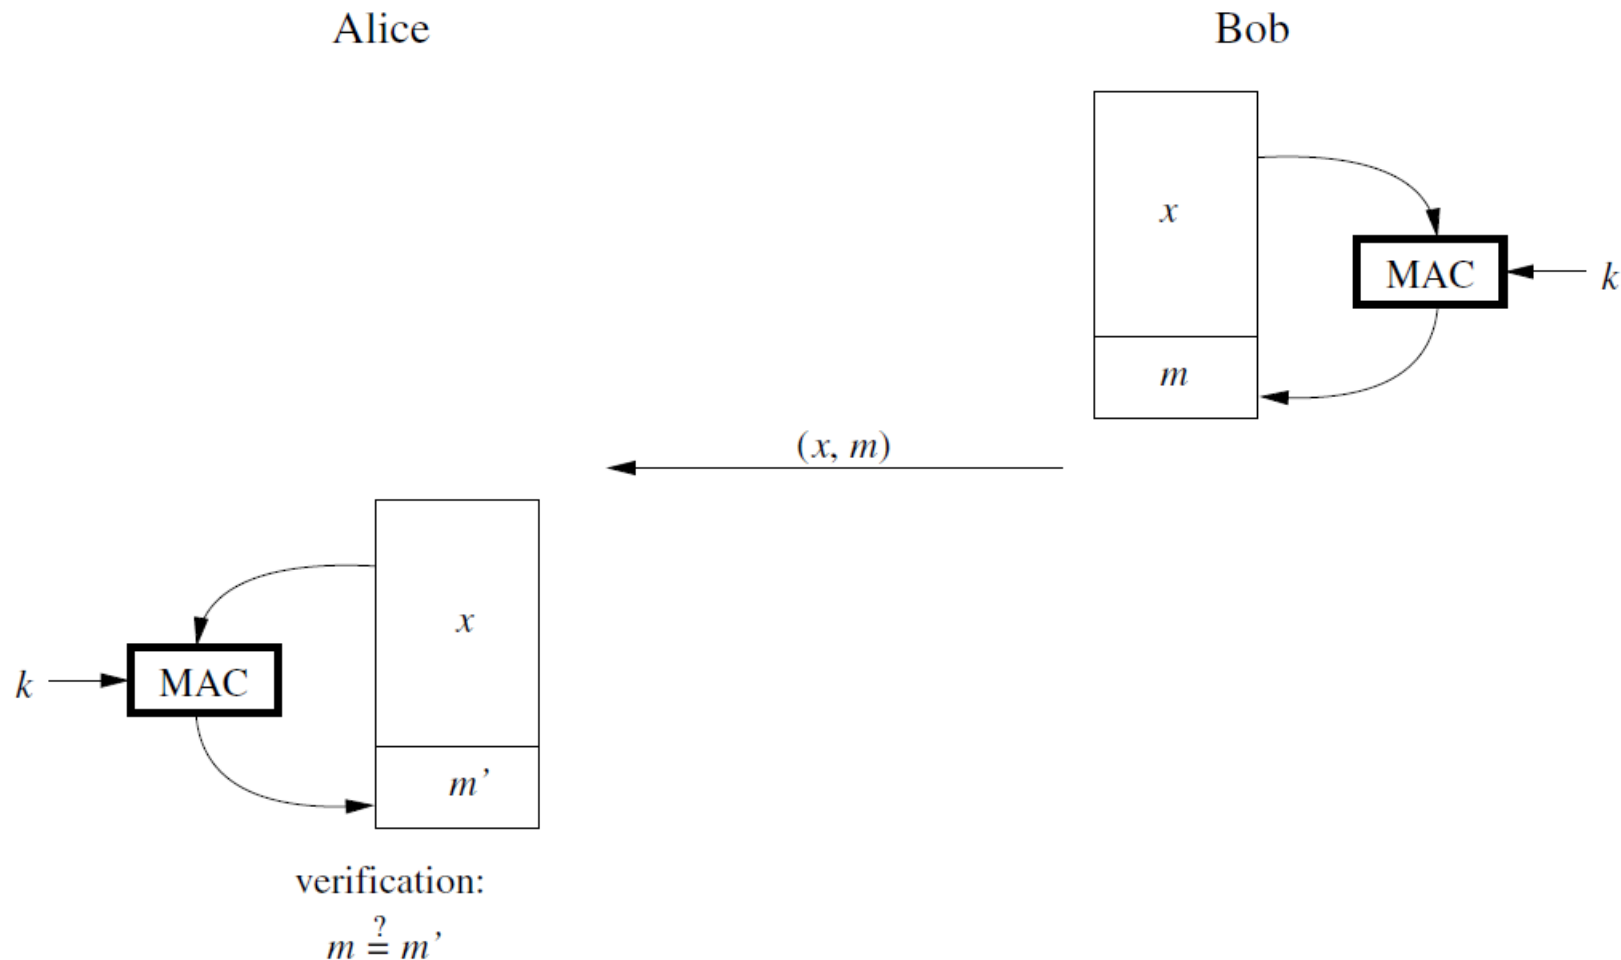
- But note that merely using a hash function, like SHA-1, cannot completely guarantee integrity
  - Eve/Oscar can still modify the message via the length extension attack
- An even worse situation: Eve/Oscar can perform a *man-in-the-middle* attack, by intercepting Alice's transmission of her plaintext-hash pair to Bob, then substituting that plaintext-hash pair with any arbitrary plaintext-hash pair of Eve/Oscar's choice – Bob has no way of knowing that the plaintext-hash pair did not come from Alice

# Message authentication codes (MACs)

- Using a *shared secret key*, combined with a hash function, can provide a better guarantee of message integrity
  - This is the basic idea of a MAC
  - MACs can also be derived from block ciphers
  - MACs offer better protection against man-in-the-middle attacks, since only Alice and Bob possess the key, and thus are the only two parties capable of computing the MAC

# Message authentication codes (MACs)

- Basic principle of MACs:



# Message authentication codes (MACs)

- Motivation for using MACs: Alice and Bob need assurance that any modifications of a plaintext message during transit are detected
- Suppose Bob wants to send Alice a plaintext message  $x$ . They both share a secret key  $k$
- He computes the authentication tag  $m$  as a function of the message  $x$  and the shared secret key  $k$ :

$$m = \text{MAC}_k(x)$$

- He then sends both the message  $x$  and the authentication tag  $m$
- When Alice receives  $x$  and  $m$ , she recomputes the authentication tag using the received message  $x$  and her key  $k$  (call this tag  $m'$ )

# Message authentication codes (MACs)

- She can then check that  $m' = m$ . If this is the case, she can be confident that the message  $x$  was unaltered during transit, since her MAC computation  $m'$  would yield a different result if  $x$  was altered
- *Message integrity* is provided as a security service by MACs, just like hash functions



# Message authentication codes (MACs)

- MACs also provide *message authentication*, in addition to message integrity
  - Alice can be assured that Bob is the source of a message and vice versa
  - This is possible because only Alice and Bob possess the secret key  $k$  and therefore are the only ones capable of computing the tag
  - Eve and Oscar are unable to compute the tag as they do not have the key
- However, MACs **cannot** provide *non-repudiation*
  - This is because the secret key  $k$  is shared by Alice and Bob, and there is no way to prove to a *neutral third party* that a *message* and its authentication tag (i.e. MAC) *originated from either Alice or Bob*

# MACs from hash functions (HMAC)

- As mentioned earlier, we can use cryptographic hash functions (like SHA-1), together with a secret key, to realize MACs
- One possible construction, HMAC, is popular and commonly used
- HMAC is used in the TLS protocol
- HMAC is secure, provided that certain assumptions are made
- Two basic ways to construct HMAC:  
 $m = MAC_k(x) = h(k \parallel x) \rightarrow$  secret **prefix** MAC  
 $m = MAC_k(x) = h(x \parallel k) \rightarrow$  secret **suffix** MAC
- “ $\parallel$ ” is the **concatenation** operator

# Attacks against HMACs

- Intuitively, both ways should result in strong cryptographic MACs, since modern cryptographic hash functions have the preimage resistance (one-wayness) property and good scrambling operations
- However, both approaches have weaknesses; let's discuss possible attacks against the secret prefix MAC and secret suffix MAC
- For these attacks, assume that the hash function is based on an MD construction, i.e. the intermediate hash value (output) of the  $i$ th iteration is  $h_i = f(h_{i-1}, x_i)$ , where  $f$  is some compression function

# Attacks against secret prefix MACs

- $m = h(k \parallel x)$
- Suppose Bob wishes to send a message  $x = (x_1, \dots, x_L)$  to Alice
- He computes an authentication tag  $m$  using his secret key  $k$ , with
$$m = \text{MAC}_k(x) = h(k \parallel x_1, \dots, x_L) = h_L = f(h_{L-1}, x_L)$$
- He then sends  $x$  and  $m$  over to Alice

length ext attack for secret prefix MACs.

$$m = \text{MAC}_k(x) = h(x \parallel k)$$

Assume  $h_i = f(h_{i-1}, x_i)$  ← old construction.

$$\text{Bob: } x = (x_1, x_2, \dots, x_L)$$

$$m = \text{MAC}(k \parallel x_1, x_2, \dots, x_L)$$

Oscar: intercept Bob's message  $x$  &  $m$ .

⇒ note: oscar cannot send any arbitrary  $x'$  &  $m'$

but he can construct  $x_{\text{oscar}} = (x_1, x_2, \dots, x_L, x_{L+1})$

$$\text{construct } m_{\text{oscar}} = h(k \parallel x_1, x_2, \dots, x_L, x_{L+1})$$

$$= f(h_L, x_{L+1})$$

$$= f(m, x_{L+1})$$

then sends  $x_{\text{oscar}}, m_{\text{oscar}}$  to Alice

extra block  
for length  
ext attack.

no 'k' involved  
here as m was  
provided by Bob.

Alice: recomputes  $m_{\text{oscar}}$  using  $x_{\text{oscar}}$

⇒ valid tag as  $m_{\text{oscar}} = m'$

⇒ thinks  $x_{\text{oscar}}$  is authentic

# Attacks against secret prefix MACs

- However, Oscar intercepts the transmission containing  $x$  and  $m$ . He constructs an altered message  $x_{Os} = (x_1, \dots, x_L, x_{L+1})$  by adding an additional block  $x_{L+1}$  to the original message
- He then computes an authentication tag  $m_{Os}$ , with
$$m_{Os} = h(k \parallel x_1, \dots, x_L, x_{L+1}) = h_{L+1} = f(h_L, x_{L+1}) = f(m, x_{L+1})$$
- In other words,  $m_{Os} = f(m, x_{L+1}) = h(m \parallel x_{L+1})$ , which implies that Oscar does **not** need to know the secret key  $k$  to compute  $m_{Os}$

# Attacks against secret prefix MACs

- Oscar sends  $x_{Os}$  and  $m_{Os}$  over to Alice, who computes  $m' = MAC_k(x_{Os}) = h(k || x_{Os}) = h(k || x_1, \dots, x_L, x_{L+1}) = m_{Os}$
- Alice accepts the message  $x_{Os}$  as authentic, since  $m' = m_{Os}$ 
  - Note that  $m \neq m_{Os}$ , but this does not matter as  $m_{Os}$  itself is a **valid** tag

# Attacks against secret suffix MACs

- $m = h(x \parallel k)$
- Again, suppose Bob wishes to send a message  $x$  to Alice
- He computes an authentication tag  $m$  using his secret key  $k$ , with  
 $m = MAC_k(x) = h(x \parallel k)$
- He then sends  $x$  and  $m$  over to Alice

Bob:  $x$   
 $m = MAC_k(x) = h(x \parallel k)$

Oscar: intercepts  $x, m$   
needs to find a 2nd preimage  $x_{oscar}$   
such that  $h(x_{oscar}) = h(x)$   
⇒ challenging, but if he is successful,  
he can send  $x_{oscar}, m$  to Alice

*harder than finding collision.*

# Attacks against secret suffix MACs

- Again, Oscar intercepts the transmission containing  $x$  and  $m$
- This time, Oscar tries to create a second message  $x_{Os}$  such that  $h(x_{Os}) = h(x)$ , with  $x_{Os} \neq x$
- This is basically finding a second preimage of the hash function
- If he can find such a message  $x_{Os}$ , then  $m = h(x \parallel k) = h(x_{Os} \parallel k)$ , because of the iterative nature of the hash function
- So  $m$  would also be a valid tag for  $x_{Os}$
- Then Oscar can send  $x_{Os}$  and  $m$  over to Alice



# Attacks against HMACs: comments

- Attacks against secret prefix MACs are basically length extension attacks
- Attacks against secret suffix MACs require finding a second preimage of the hash function, which is much harder to achieve, as cryptographic hash functions are designed to have second preimage resistance
- So secret suffix MACs are better than secret prefix MACs

# Storage of secrets

- We can also use **hashes** for storage of **secret data**
- Storage of usernames and the corresponding passwords on a computer:
  - It's not a good idea to store the passwords in plaintext
  - To protect the passwords, the passwords are hashed using a cryptographic hash function
  - E.g. The username-password pair "**alice p@ssw0rd**" may be stored as "**alice 0x234a123d456efeed**" on the computer, after the password is passed through a cryptographic hash function
  - When the user Alice **inputs her password on the computer**, the entered password is **hashed** and the **resulting hash** is then **compared against the stored hash value**

# Storage of secrets

- On Linux distributions, the passwords are stored as SHA-512 hashes in the directory `/etc/shadow`
- Hashing the passwords before storage helps to maintain the secrecy of the passwords
- But there are ways to attack this storage scheme
  - The strategies of these attacks mainly revolve around finding the preimages of these hashes

# Finding the preimages of hashes

- Suppose Oscar manages to steal a list of hashed password values
- Is there a way for him to derive the original plaintext passwords?
- Yes, Oscar can create a list of commonly used passwords in plaintext, then compute the corresponding hash for each password
- He can then compare each computed hash with the hashes in the stolen list; matching hashes indicate the corresponding plaintext password
- To save time, he can use a precomputed dictionary of password hashes and their corresponding plaintext passwords, or rainbow tables
  - E.g. password cracking programs like Hashcat and John the Ripper can utilize precomputed dictionaries and/or rainbow tables

# Finding the preimages of hashes: mitigation

- One way in which we can mitigate such attacks on password hashes is to salt the plaintext password before hashing it
- Add some additional bits of data to the plaintext password
  - These bits of data are usually derived from the corresponding username of that password
  - Alternatively, the bits may be derived from a pseudo-random number generator
- Makes dictionaries and rainbow tables significantly less effective, because the same password for different usernames will map to different hash values
- Also increases the input space of the hash function

# Other examples of attacks on hashes

- Yuval's square root attack
- MD5 Collisions Inc.