

50.042 FCS Summer 2024

Lectures 7 & 8 – Block Ciphers

Felix LOH

Singapore University of Technology and Design



SINGAPORE UNIVERSITY OF
TECHNOLOGY AND DESIGN

With selected materials adapted from: *Understanding Cryptography: A Textbook for Students and Practitioners*, by C. Paar and J. Pelzl

Recap of basic ciphers

- Simple ciphers:
 - E.g. shift ciphers like Caesar's cipher
 - Problems: Small key space; vulnerable to frequency analysis; ciphers are linear
- Stream ciphers:
 - E.g. one-time pad (OTP)
 - Problems: Generation of the key stream can be impractical; low throughput in general

Block ciphers: motivation

- Simple ciphers have a small keyspace and are vulnerable to frequency analysis
- Stream ciphers generally suffer from low throughput and have issues with the key stream generation
- So we need more efficient ciphers/cryptoalgorithms that can encrypt large data sets
 - Parallelize part of the operations of the cipher
 - Leverage the large word width of modern CPUs

Primitive operations for block ciphers


- There are two primitive operations with which strong encryption algorithms can be built
- **Confusion:** An encryption operation where the relationship between the key and ciphertext is obscured
 - Can be achieved with the substitution operation
- **Diffusion:** An encryption operation where the influence of one plaintext symbol is spread over many ciphertext symbols with the goal of hiding the statistical properties of the plaintext
 - Can be achieved with bit permutation operations
 - Modern block ciphers possess excellent diffusion properties – this means that changing one bit of the plaintext results, on average, in the change of half of the bits of the ciphertext

→ spread the influence of one plaintext bit over many ciphertext bits

eg. $x_1 = 0010$ $x_2 = 0000$ $\begin{matrix} 1011 \\ 1011 \end{matrix}$ → block cipher → $y_1 = 1011$ $y_2 = 0110$ $\left. \begin{matrix} 1001 \\ 1100 \end{matrix} \right\}$ many difference


Primitive operations for block ciphers

- **Both** confusion and diffusion operations are utilized in modern block ciphers
- Several iterations of one diffusion operation followed by a confusion operation
- This is also called a product cipher
 - In other words, all modern block ciphers are product ciphers



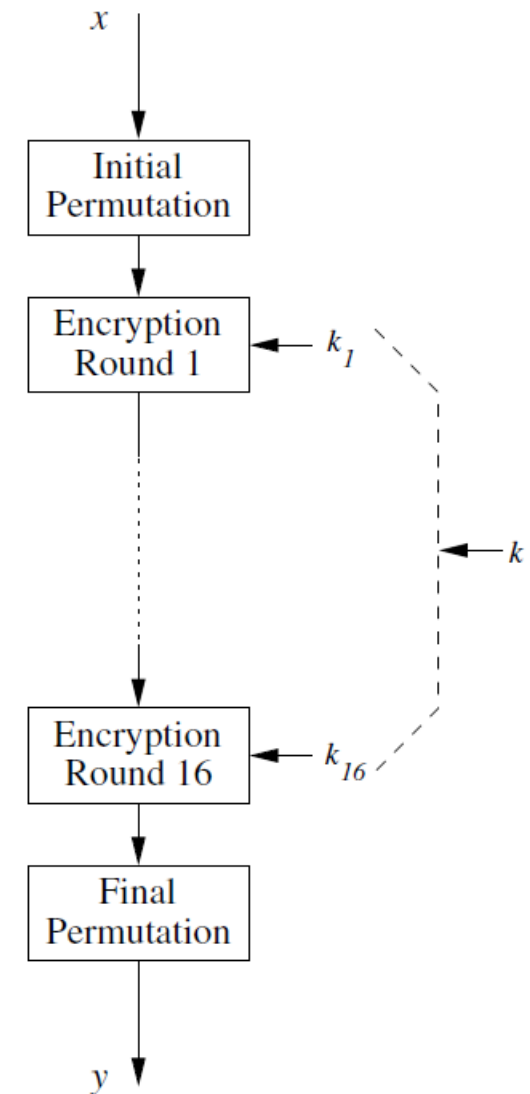
Overview of the DES algorithm

- DES: **Data Encryption Standard**, a popular cipher released in 1977, developed by IBM with the assistance of the NSA
- Encrypts 64-bit blocks of plaintext, using a 56-bit key
 - DES originally had a 128-bit key, but apparently the NSA convinced IBM to reduce the key length to 56 bits
 - This made DES more vulnerable to brute-force attacks
- DES is a symmetric cipher, i.e. the same key is used for both encryption and decryption
- DES was the cipher of choice of the US government, until it was replaced by the AES cipher in 1999
 - Variants of DES are still in use, e.g. 3DES is commonly used in smart cards and credit cards



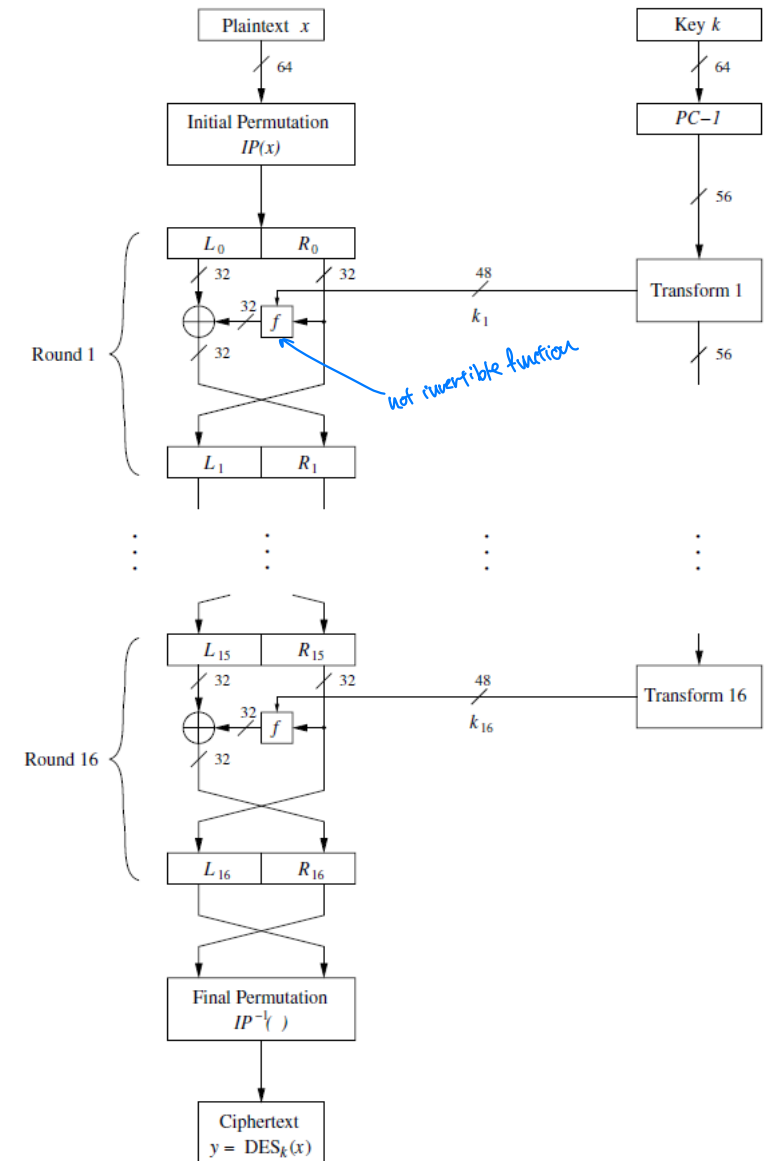
Overview of the DES algorithm

- The encryption of each 64-bit block of plaintext is handled in 16 rounds
- Each round performs the same operation
 - A different round key is used in each round; round key k_i is used in Round i — e.g. round key k_1 is used for Round 1, and so on
 - All round keys k_i are derived from the main key k (this round key derivation is known as a key schedule)
- Decryption is carried out in a similar manner; round key k_{16} is used for Round 1 (a decryption round is the same as an encryption round), round key k_{15} is used for Round 2, and so on, with round key k_1 used for Round 16



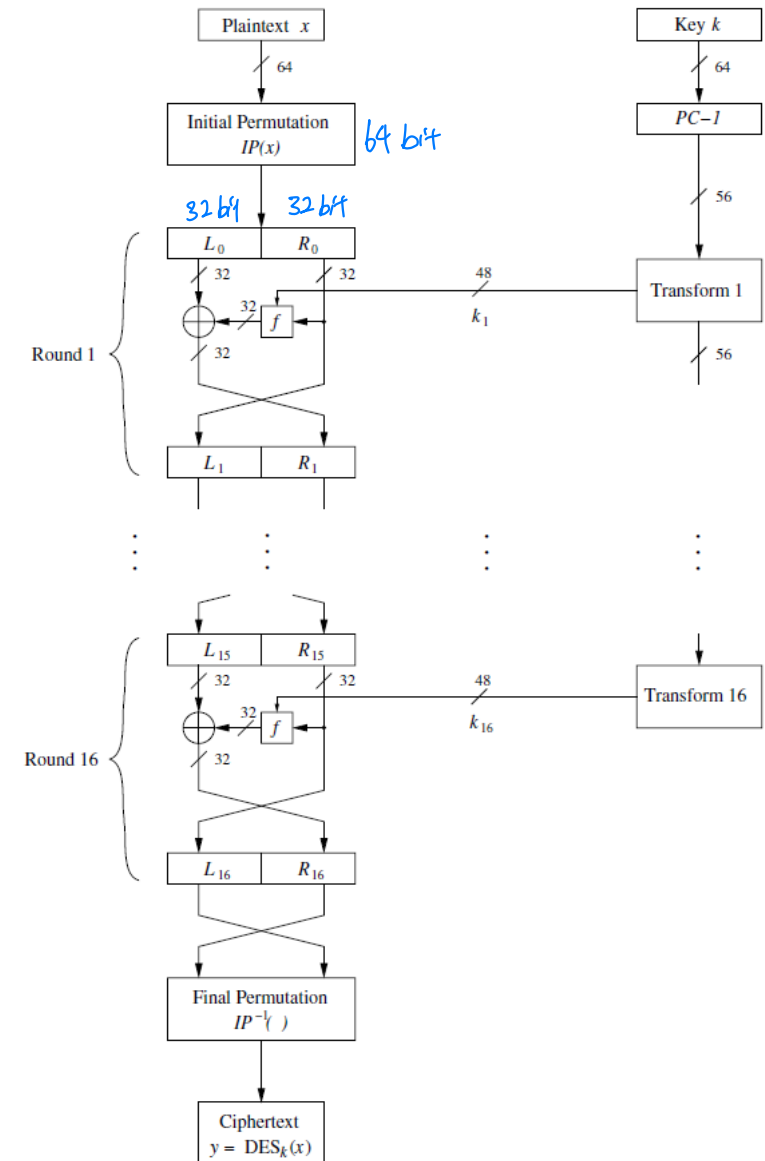
The DES algorithm: encryption

- The structure shown on the right is called a *Feistel network*
 - If designed carefully, Feistel networks can result in strong ciphers
 - One advantage of Feistel networks: encryption and decryption are *almost the same process* – decryption only requires a reversed key schedule
 - Feistel networks are used in *many (but not all)* modern *block ciphers*
- The key k is *listed as 64 bits*; this is because 8 of those bits are used as *parity bits* and are *not actually key bits* – thus the key size is *56 bits*



The DES algorithm: encryption

1. There is an initial bitwise permutation IP of the 64-bit plaintext x
2. The plaintext is then split into two 32-bit halves L_0 and R_0 , which are fed into the Feistel network, which consists of 16 rounds
3. The right half R_i is used as an input to the f -function, whose output is then XOR-ed with the left half L_i
4. The left and right halves are then swapped



The DES algorithm: encryption

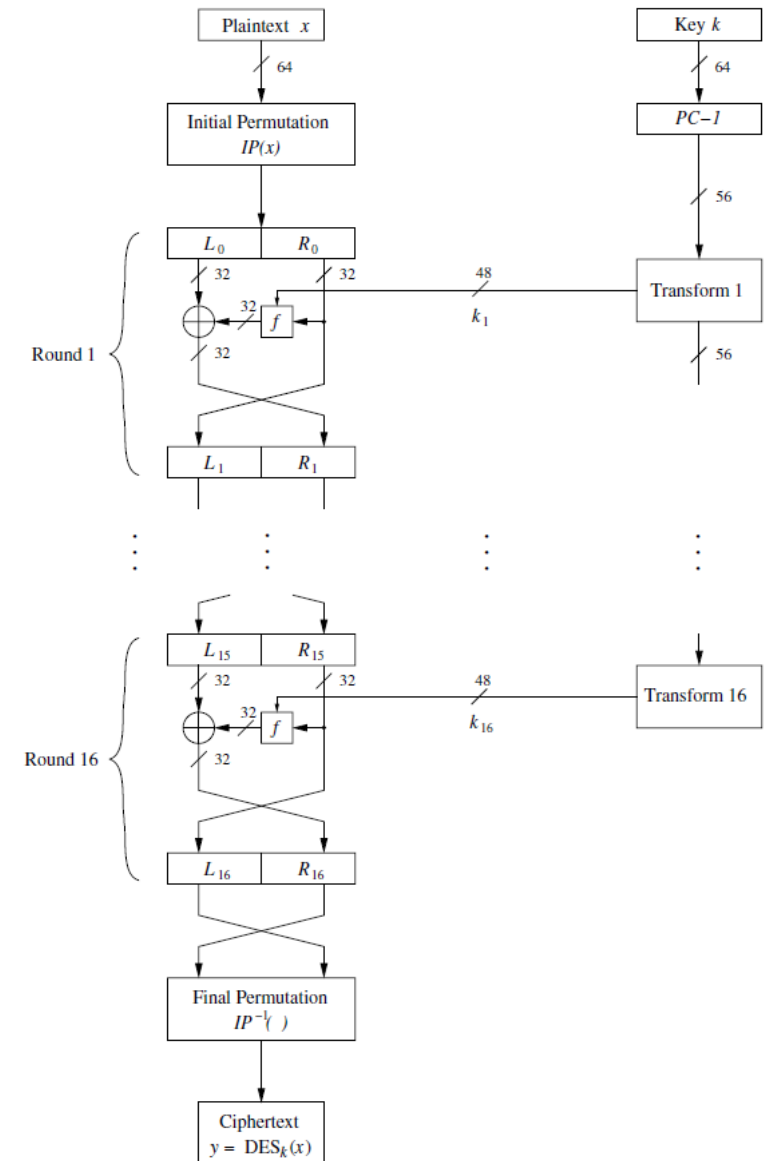
5. Steps 3 and 4 are repeated in the next round, and the operations in each round can be expressed as:

Wikipedia


$$L_i = R_{i-1}$$

$$R_i = L_{i-1} \oplus f(R_{i-1}, k_i), \text{ where } i = 1, \dots, 16$$


6. After Round 16, the 32-bit halves L_{16} and R_{16} are swapped again
7. There is a final bitwise permutation IP^{-1} of the resulting 64-bit value (IP^{-1} is the inverse of IP); this results in the 64-bit ciphertext output



The DES algorithm: initial and final permutation



IP							
58	50	42	34	26	18	10	2
60	52	44	36	28	20	12	4
62	54	46	38	30	22	14	6
64	56	48	40	32	24	16	8
57	49	41	33	25	17	9	1
59	51	43	35	27	19	11	3
61	53	45	37	29	21	13	5
63	55	47	39	31	23	15	7

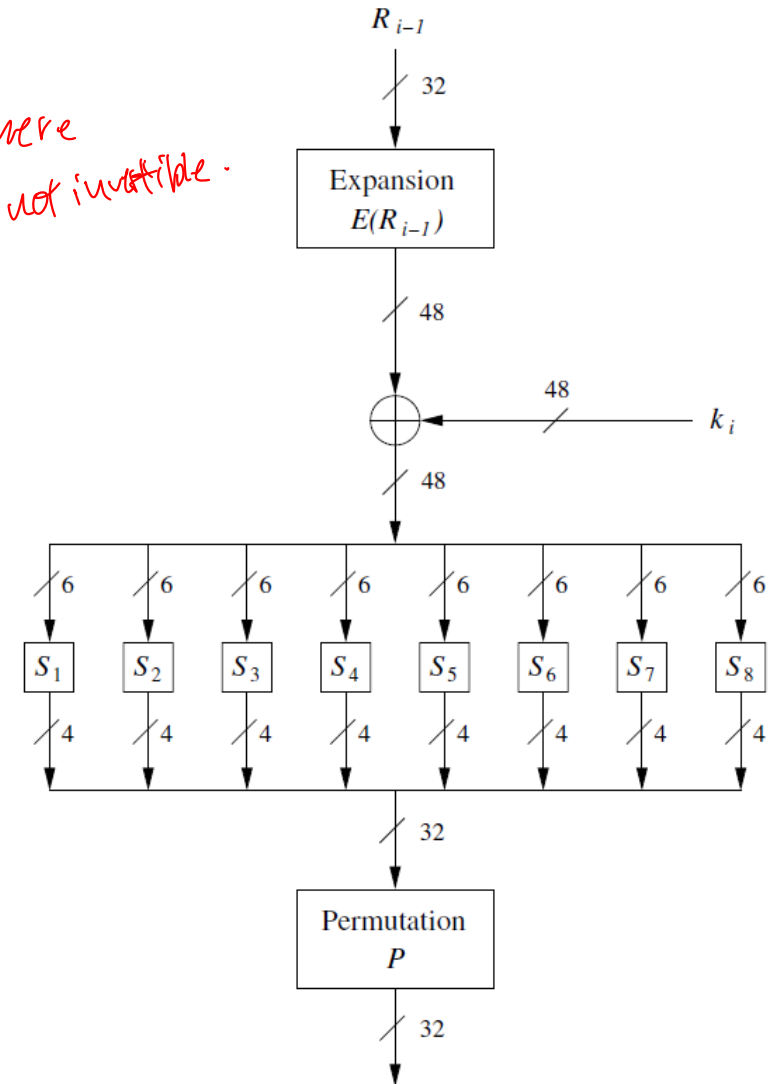


IP^{-1}							
40	8	48	16	56	24	64	32
39	7	47	15	55	23	63	31
38	6	46	14	54	22	62	30
37	5	45	13	53	21	61	29
36	4	44	12	52	20	60	28
35	3	43	11	51	19	59	27
34	2	42	10	50	18	58	26
33	1	41	9	49	17	57	25

- IP : Bit position 1 of plaintext x is mapped to bit position 40; position 58 mapped to 1, etc.
- IP^{-1} : Bit position 40 of z (output of Round 16) is mapped to bit position 1 (inverse of IP)

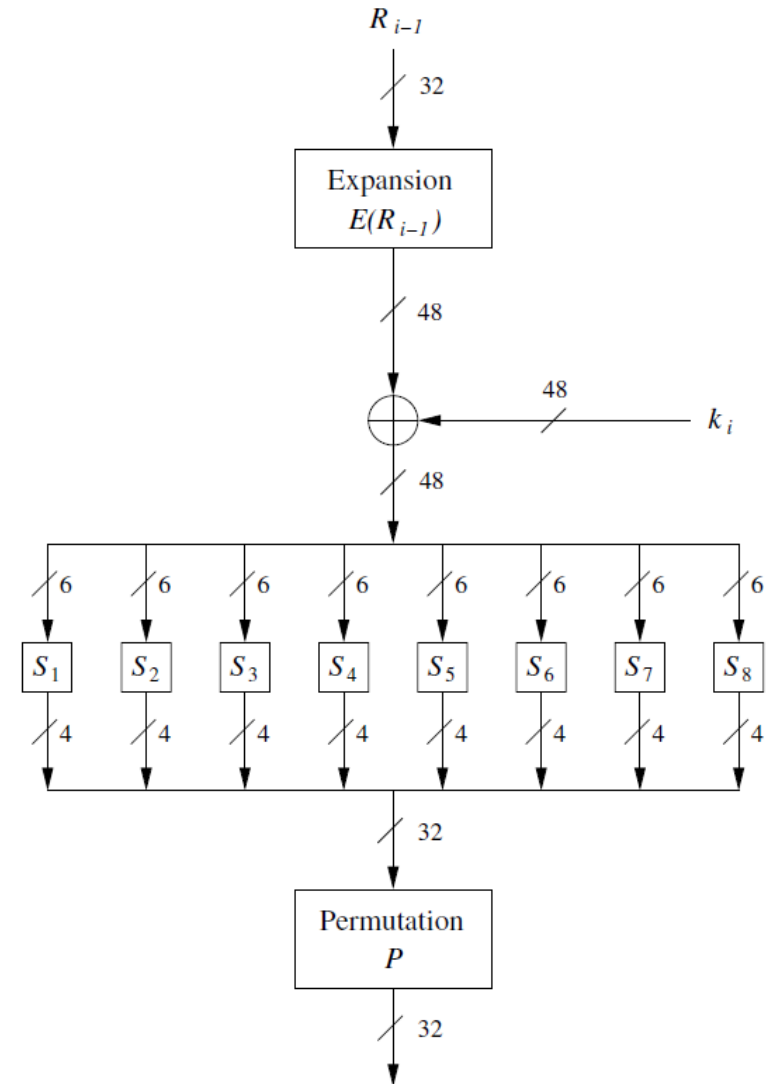
The DES algorithm: f -function

1. The 32-bit R_{i-1} input is expanded to 48 bits
 - 16 of the 32 input bits are copied \rightarrow information loss here
 \downarrow
hence f not invertible.
 - Some permutation is also done here
2. The 48-bit result of the expansion is XOR-ed with the round key k_i , then separated into eight 6-bit blocks
3. Each 6-bit block is operated on by a different substitution box (S-box), resulting in a 4-bit value
4. The 4-bit blocks are recombined into a 32-bit value and then permuted bitwise using a permutation operation P



The DES algorithm: f -function

- The two primitive operations, **confusion** and **diffusion**, are realized within the f -function
- The S -boxes provide confusion
 - These S -boxes are **non-linear**, i.e.
 $S(a) \oplus S(b) \neq S(a \oplus b)$
- The permutation operation P provides diffusion
 - This is because the 4-bit outputs of the S -boxes are **permuted** such that they **affect several different S -boxes** in the **following round**




The DES algorithm: S-box example

- An S-box is basically a lookup table which maps a 6-bit input to a 4-bit output
- E.g. S-box S_1 is as follows:

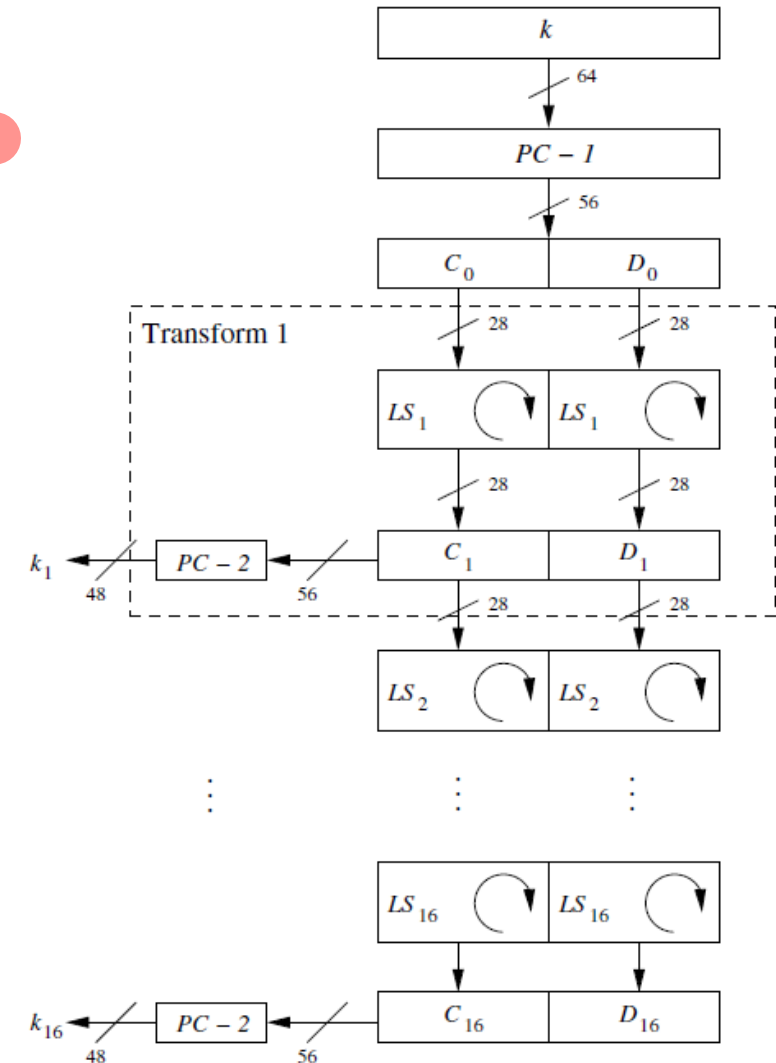
S_1	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	14	04	13	01	02	15	11	08	03	10	06	12	05	09	00	07
1	00	15	07	04	14	02	13	01	10	06	12	11	09	05	03	08
2	04	01	14	08	13	06	02	11	15	12	09	07	03	10	05	00
3	15	12	08	02	04	09	01	07	05	11	03	14	10	00	06	13

- The MSB and LSB of the 6-bit input select the row; the 4 inner bits select the column
 - Suppose we have a 6-bit input 100101_2 to S_1
 - Choose the 4th row (3_{10}) and 3rd column (2_{10})
 - So the 4-bit output is $08_{10} = 1000_2$



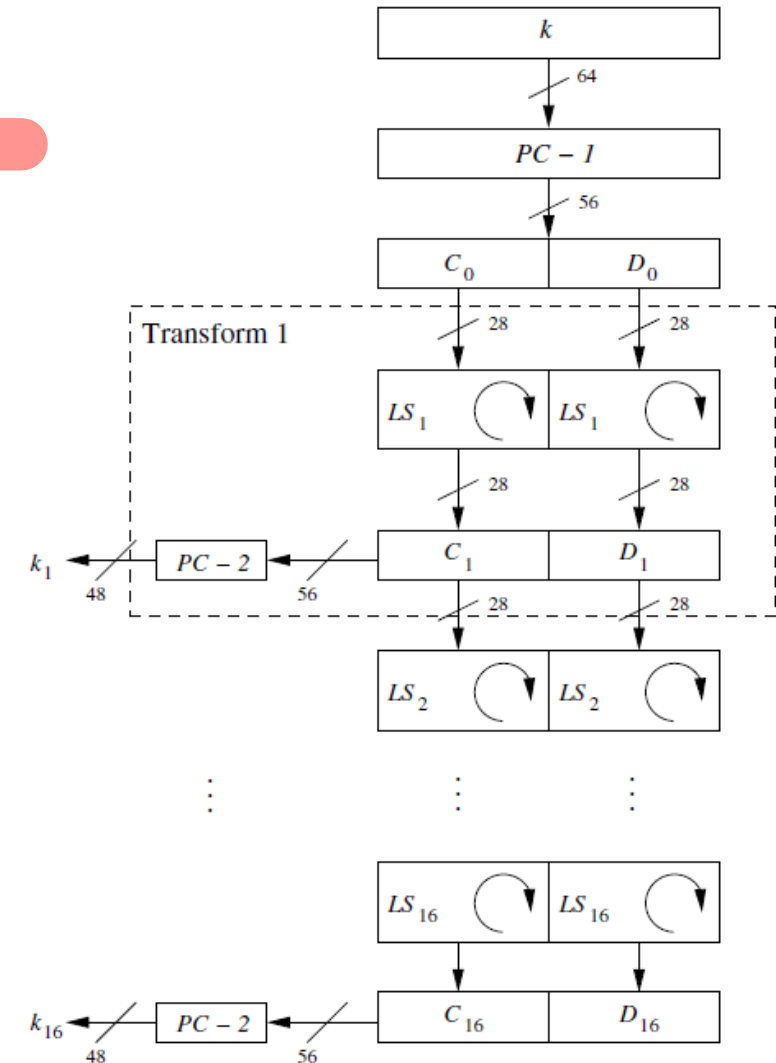
The DES algorithm: key schedule

- The key schedule **derives** 16 round keys k_i , each consisting of **48 bits from the original 56-bit key**
1. During the $PC-1$ operation, the 64-bit key (including the 8 parity bits) is **reduced to 56 bits** by **discarding every eighth bit** (the parity bit) and the remaining bits are **permuted**
 2. The resulting 56-bit value is then split into **28-bit halves C_{i-1} and D_{i-1}** , and are passed on to the ***Transform i* operation**, for $i = 1, \dots, 16$




The DES algorithm: key schedule

- Next, during the 1st phase of the *Transform i* operation, the two halves are each circular left shifted, by one or two bits depending on the round *i*, resulting in C_i and D_i
- During the 2nd phase of the *Transform i* operation, the 56 bits from the two halves C_i and D_i are permuted using the *PC-2* operation (also dropping 8 bits), resulting in round key k_i
- Steps 3 and 4 are repeated for the derivation of the next round key, using the two halves C_i and D_i from the 1st phase of *Transform i*



The DES algorithm: decryption

- Decryption is carried out in a similar manner; there is an initial bitwise permutation IP of the ciphertext
- Then subkey k_{16} is used for round 1, round key k_{15} is used for round 2, and so on, until round key k_1 is used for round 16
- This effectively reverses the encryption in a round-by-round manner
 - Round 1 of decryption reverses round 16 of encryption, round 2 of decryption reverses round 15 of encryption, and so on until round 16 of decryption reverses round 1 of encryption
- The final bitwise permutation IP^{-1} of decryption reverses the initial bitwise permutation IP of encryption; this results in the original plaintext x



The DES algorithm: decryption

- Let's show that the decryption process reverses the encryption in a **round-by round manner**

- Note that the **f -function itself is not invertible**

Let $x = (L_0, R_0) \leftarrow$ one 64-bit block
 y is the ciphertext

after permutation,
of y

$$(L_0^d, R_0^d) = IP(y) = IP(IP^{-1}(R_{16}, L_{16}))$$

$$= (R_{16}, L_{16})$$

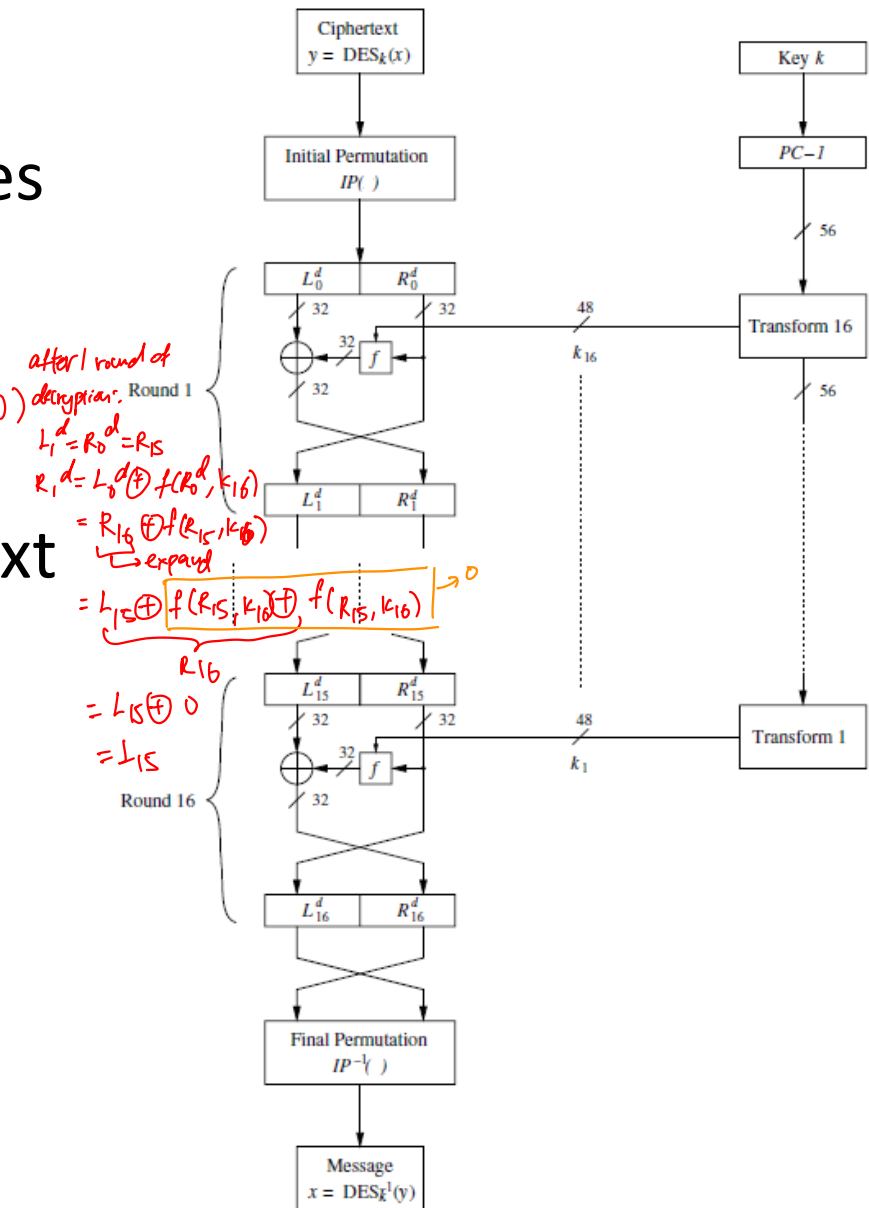
$$\Rightarrow L_0^d = R_{16}, R_0^d = L_{16} = R_{15}$$

- After the initial permutation IP of the ciphertext y , we have

$$(L_0^d, R_0^d) = IP(y) = IP(IP^{-1}(R_{16}, L_{16})) = (R_{16}, L_{16})$$

- So, we have

$$L_0^d = R_{16} \text{ and } R_0^d = L_{16} = R_{15}$$




The DES algorithm: decryption

- Now the output (L_1^d, R_1^d) of round 1 of decryption can be expressed as follows:

$$L_1^d = R_0^d = R_{15}$$

$$\begin{aligned} R_1^d &= L_0^d \oplus f(R_0^d, k_{16}) \\ &= R_{16} \oplus f(R_{15}, k_{16}) \\ &= L_{15} \oplus f(R_{15}, k_{16}) \oplus f(R_{15}, k_{16}) \\ &= L_{15} \oplus 0 \\ &= L_{15} \end{aligned}$$

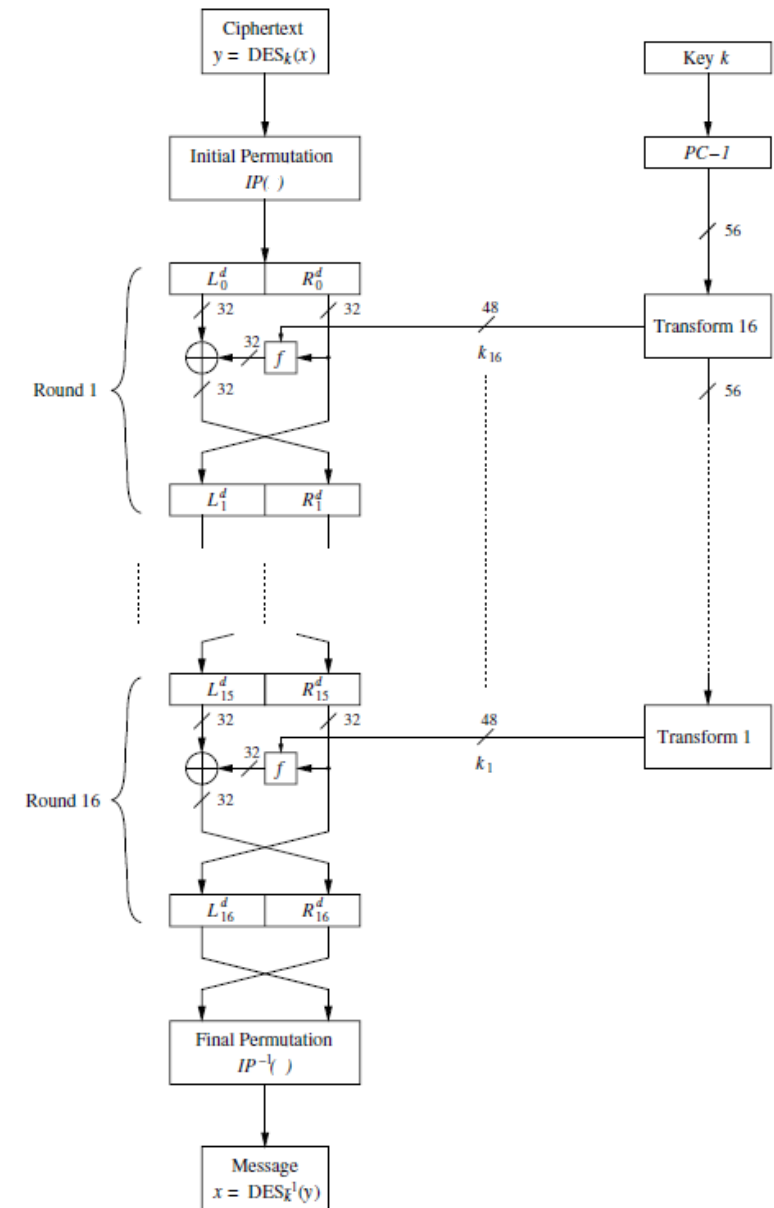


The DES algorithm: decryption

- Thus, the output of round 1 of decryption is $(L_1^d, R_1^d) = (R_{15}, L_{15})$
- R_{15} and L_{15} are effectively the input bits to round 16 of encryption; we have shown that round 1 of decryption reverses round 16 of encryption
- So we can express the decryption process iteratively as:

$$L_i^d = R_{16-i}$$

$$R_i^d = L_{16-i}, \text{ where } i = 0, 1, \dots, 16$$



The DES algorithm: decryption


- Particularly, after round 16 of decryption, we have:

$$L_{16}^d = R_0$$

$$R_{16}^d = L_0$$

- Thus, after the final bitwise permutation IP^{-1} , we have

$$\begin{aligned} IP^{-1}(R_{16}^d, L_{16}^d) &= IP^{-1}(L_0, R_0) \\ &= IP^{-1}(IP(x)) \\ &= x \end{aligned}$$




The DES algorithm: security analysis

- The DES block cipher is now vulnerable to brute-force attacks (i.e. exhaustive key search) due to the small key length of 56 bits
 - Keyspace is too small for today's security requirements
 - Given at least one pair of plaintext-ciphertext (x, y) , an attacker can test all possible keys, by checking whether $DES_{k_i}^{-1}(y) = x$, for $i = 0, 1, 2, 3, \dots, 2^{56} - 1$
- No practical analytical attacks found
- Use the 3DES variant as a replacement for DES

much less than
 2^{80}
 2^{56}

3DES algorithm



- DES is no longer secure, but its variant 3DES is still viable and in current use
- 3DES uses three 56-bit keys k_1, k_2, k_3 and computes the 64-bit block of ciphertext y as follows:

$$y = DES_{k_3}(DES_{k_2}^{-1}(DES_{k_1}(x)))$$


- If $k_1 = k_2 = k_3$, then 3DES effectively behaves like DES (allows for backward compatibility with legacy systems)

3DES algorithm

- 3DES is commonly used in credit cards and smart cards, for authentication of transactions
 - The 3DES block cipher is being utilized as a MAC in these situations
- 3DES increases the effective key length to 112 bits (as opposed to the 56-bit key length of DES)
 - Improved security over DES
 - Note: The effective key length of 3DES is **not** increased to 168 bits, because of the *meet-in-the-middle* attack

Overview of the AES algorithm

- AES: **A**dvanced **E**ncryption **S**tandard, introduced in 2001
 - This cipher is the result of an NIST call for proposals for a new cipher to replace DES
- Encrypts 128-bit blocks of plaintext, with a choice of key lengths of 128, 192 or 256 bits
 - To date, there are no known attacks better than brute-force attacks against AES
- AES is a symmetric cipher, i.e. the same key k is used for both encryption and decryption
 - Unlike DES, AES does not utilize a Feistel network
- AES is the current cipher of choice of the US government
 - Used by US federal agencies to encrypt classified documents
- Also used in protocols like IPsec, TLS, SSH



Overview of the AES algorithm

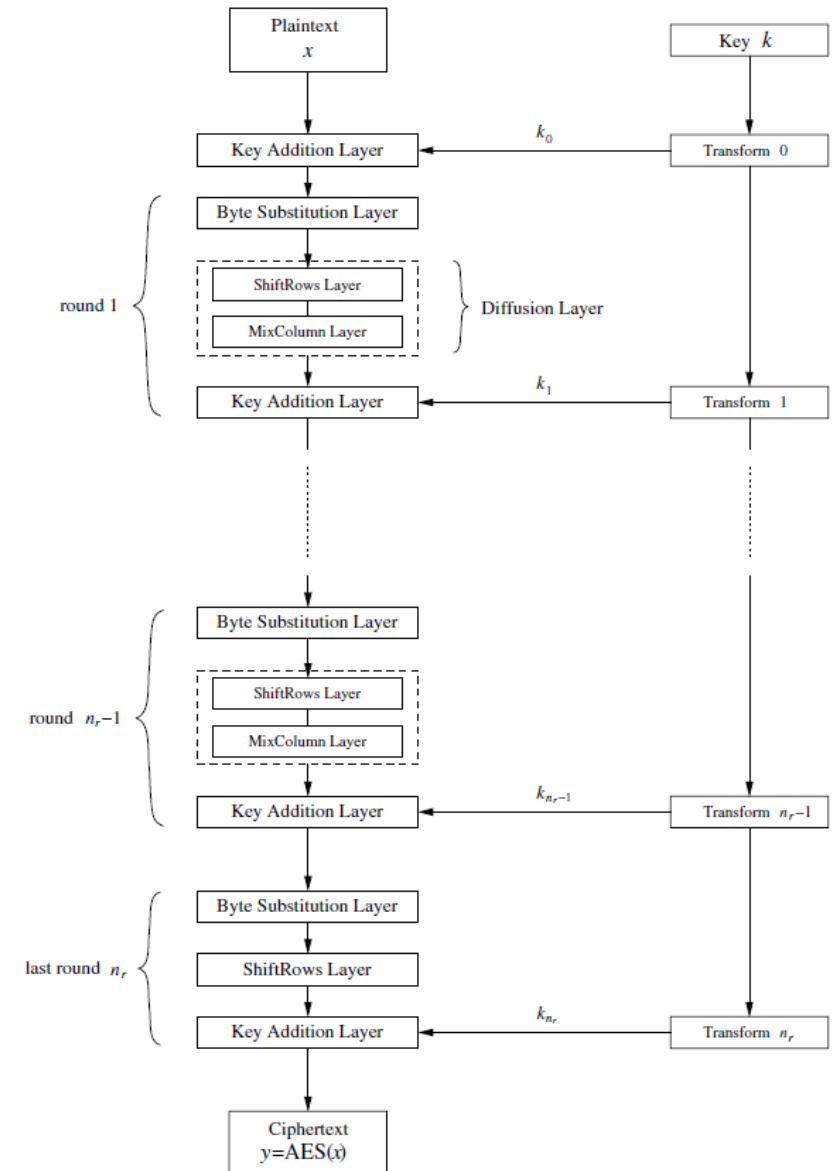
- The encryption of each 128-bit block of plaintext is handled in several rounds; the number of rounds needed, n_r , depends on the key length:
 - 128-bit key length — $n_r = 10$
 - 192-bit key length — $n_r = 12$
 - 256-bit key length — $n_r = 14$
- Each round of AES encrypts all 128 bits of the data
 - By contrast, DES only encrypts 32 bits of the 64-bit data block in each round
- Similar to DES, each round of AES uses a round key k_i which is derived from the key k using a key schedule
- AES encryption consists of *layers*, each of which manipulates all 128 bits of the data path in some manner (the data path is known as the *state*)
- Decryption requires a reversed key schedule, and all layers utilized during the encryption process need to be inverted

The AES algorithm: encryption

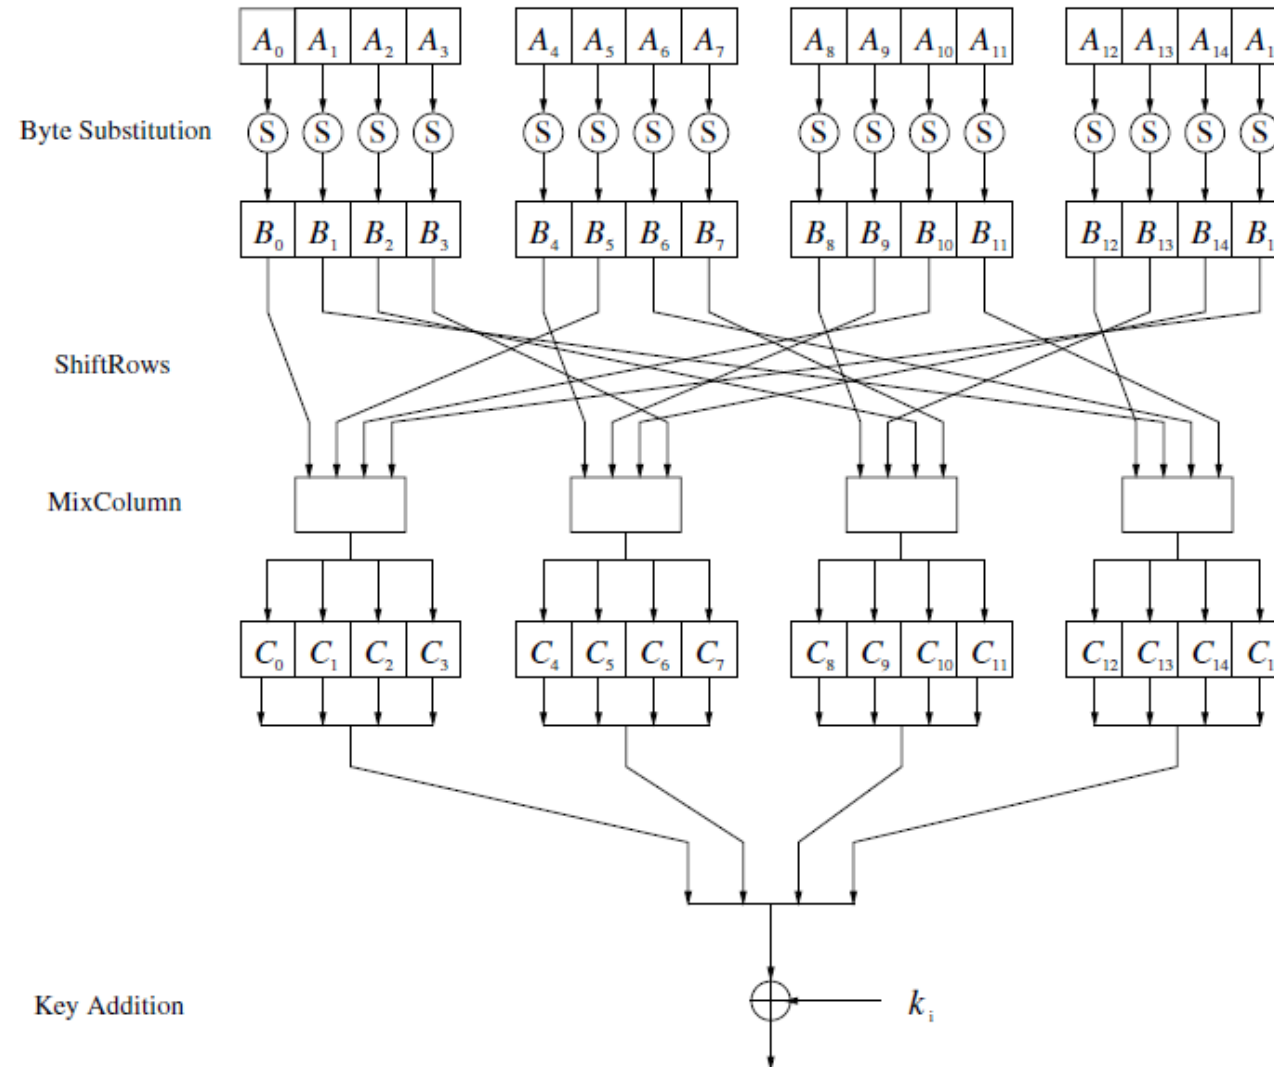
- A brief description of the three layers utilized in AES encryption is as follows
- **Key Addition layer:** A 128-bit round key is XOR-ed to the state
- **Byte Substitution layer (S-box):** Each element of the state is non-linearly transformed using lookup tables with special mathematical properties
 - This introduces confusion to the data
- **Diffusion layer:** Provides diffusion over all state bits, and it consists of two sublayers (both sublayers are linear operations)
 - **ShiftRows sublayer:** Permutes the data on a byte level
 - **MixColumn sublayer:** A matrix operation which combines blocks of 4 bytes

The AES algorithm: encryption

- Each round of encryption uses all three layers, except the last round n_r , which does not utilize the MixColumn sublayer



The AES algorithm: encryption (single round)



The AES algorithm: encryption (single round)

- It is helpful to visualize the 16 byte data path (i.e. input state matrix) $A = (A_0, A_1, \dots, A_{15})$ that is moving through the AES encryption process as a 4×4 state matrix, with each of the bytes A_0, \dots, A_{15} arranged as follows:

A_0	A_4	A_8	A_{12}
A_1	A_5	A_9	A_{13}
A_2	A_6	A_{10}	A_{14}
A_3	A_7	A_{11}	A_{15}

- AES operates on elements, columns or rows of the current state matrix

The AES algorithm: encryption (single round)

1. The 16 byte (i.e. 128 bits) input state matrix $A = (A_0, A_1, \dots, A_{15})$ is fed byte-wise into the Byte Substitution layer
2. Then, the resulting 16 byte state matrix $B = (B_0, B_1, \dots, B_{15})$ is permuted byte-wise in ShiftRows sublayer
3. The result of the ShiftRows sublayer is then mixed by combining selected blocks of 4 bytes in the MixColumns sublayer
4. Lastly, the output of the ShiftRows sublayer, the state matrix $C = (C_0, C_1, \dots, C_{15})$ is XOR-ed with the round key k_i

The AES algorithm: Byte Substitution layer

- This layer is basically a row of 16 parallel S-boxes
 - Each S-box has an 8-bit input and an 8-bit output (i.e. one byte input/output)
 - All 16 S-boxes are identical, unlike DES which uses 8 different S-boxes
- Each S-box substitutes the input byte A_i with another byte B_i
 - So each S-box performs the operation $B_i = S(A_i)$
 - The substitution is a bijective mapping; each of the 2^8 possible different byte values is one-to-one mapped to another byte value
 - This allows the S-boxes to be inverted/reversed (unlike the S-boxes for DES), which is necessary for decryption
- The S-boxes are the only non-linear elements in AES
 - So $\text{ByteSub}(A) + \text{ByteSub}(B) \neq \text{ByteSub}(A + B)$ for any two arbitrary state matrices A and B

The AES algorithm: Byte Substitution layer

- S-box substitution values for an input byte (x, y):

		y															
		0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
x	0	63	7C	77	7B	F2	6B	6F	C5	30	01	67	2B	FE	D7	AB	76
	1	CA	82	C9	7D	FA	59	47	F0	AD	D4	A2	AF	9C	A4	72	C0
	2	B7	FD	93	26	36	3F	F7	CC	34	A5	E5	F1	71	D8	31	15
	3	04	C7	23	C3	18	96	05	9A	07	12	80	E2	EB	27	B2	75
	4	09	83	2C	1A	1B	6E	5A	A0	52	3B	D6	B3	29	E3	2F	84
	5	53	D1	00	ED	20	FC	B1	5B	6A	CB	BE	39	4A	4C	58	CF
	6	D0	EF	AA	FB	43	4D	33	85	45	F9	02	7F	50	3C	9F	A8
	7	51	A3	40	8F	92	9D	38	F5	BC	B6	DA	21	10	FF	F3	D2
	8	CD	0C	13	EC	5F	97	44	17	C4	A7	7E	3D	64	5D	19	73
	9	60	81	4F	DC	22	2A	90	88	46	EE	B8	14	DE	5E	0B	DB
	A	E0	32	3A	0A	49	06	24	5C	C2	D3	AC	62	91	95	E4	79
	B	E7	C8	37	6D	8D	D5	4E	A9	6C	56	F4	EA	65	7A	AE	08
	C	BA	78	25	2E	1C	A6	B4	C6	E8	DD	74	1F	4B	BD	8B	8A
	D	70	3E	B5	66	48	03	F6	0E	61	35	57	B9	86	C1	1D	9E
	E	E1	F8	98	11	69	D9	8E	94	9B	1E	87	E9	CE	55	28	DF
	F	8C	A1	89	0D	BF	E6	42	68	41	99	2D	0F	B0	54	BB	16

- E.g. If $A_0 = 0xF3$, then $B_0 = S(A_0) = S(0xF3) = 0x0D$

The AES algorithm: Byte Substitution layer

- The AES S-box implementation of one-to-one byte mappings is **not** based on some randomly chosen set of mappings (unlike DES)
- The S-box implementation is actually a two step mathematical transformation (non-linear)
 - The 1st step is a Galois field inversion for the Galois extension field $GF(2^8)$
 - The 2nd step is an affine mapping
- ***Note: We will discuss Galois fields (which includes extension fields) and operations in these fields (addition, subtraction, multiplication and inversion) during next week's lectures on modular arithmetic***

The AES algorithm: ShiftRows sublayer

- The ShiftRows sublayer is part of the Diffusion layer
 - Note: The Diffusion layer is linear, i.e. $\text{Diff}(A) + \text{Diff}(B) = \text{Diff}(A + B)$ for any two arbitrary state matrices A and B
- The ShiftRows sublayer cyclically shifts the rows of the input state matrix to the right
 - First row: no shift
 - Second row: cyclic shift by 3 bytes to the right
 - Third row: cyclic shift by 2 bytes to the right
 - Fourth row: cyclic shift by 1 byte to the right

The AES algorithm: ShiftRows sublayer

- In other words, the input state matrix $B = (B_0, B_1, \dots, B_{15})$ results in the following output:

B_0	B_4	B_8	B_{12}
B_1	B_5	B_9	B_{13}
B_2	B_6	B_{10}	B_{14}
B_3	B_7	B_{11}	B_{15}

B

→

B_0	B_4	B_8	B_{12}
B_5	B_9	B_{13}	B_1
B_{10}	B_{14}	B_2	B_6
B_{15}	B_3	B_7	B_{11}

B_{SR}

no shift

→ three positions right shift

→ two positions right shift

→ one position right shift

The AES algorithm: MixColumn sublayer

- The MixColumn sublayer is a linear transformation which mixes each column of the input state matrix
 - It is also part of the Diffusion layer
 - The MixColumn operation is the major diffusion element in AES, since every input byte influences four output bytes
- The operation performed is $C = \text{MixColumn}(B_{SR})$, where B_{SR} is the input state matrix after the ShiftRows operation is executed
- Each 4-byte column of B_{SR} is treated as a vector and multiplied by a fixed constant 4×4 matrix. Multiplication and addition of the coefficients is done in the Galois extension field $GF(2^8)$

The AES algorithm: MixColumn sublayer

- E.g. For the first column of C and B_{SR} respectively:

$$\begin{pmatrix} C_0 \\ C_1 \\ C_2 \\ C_3 \end{pmatrix} = \begin{pmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{pmatrix} \begin{pmatrix} B_0 \\ B_5 \\ B_{10} \\ B_{15} \end{pmatrix}$$

- Each state byte C_i and B_i with $i = 0, \dots, 15$, is an 8-bit value that represents an element from the Galois extension field $GF(2^8)$
- Here, column (C_0, C_1, C_2, C_3) is the result of a matrix-vector multiplication of the fixed constant 4×4 matrix with the first column of B_{SR}
 - Additions in the matrix-vector multiplication are $GF(2^8)$ additions (these are basically bitwise XOR operations)
 - Multiplications in the matrix-vector multiplication are $GF(2^8)$ multiplications

The AES algorithm: Key Addition layer

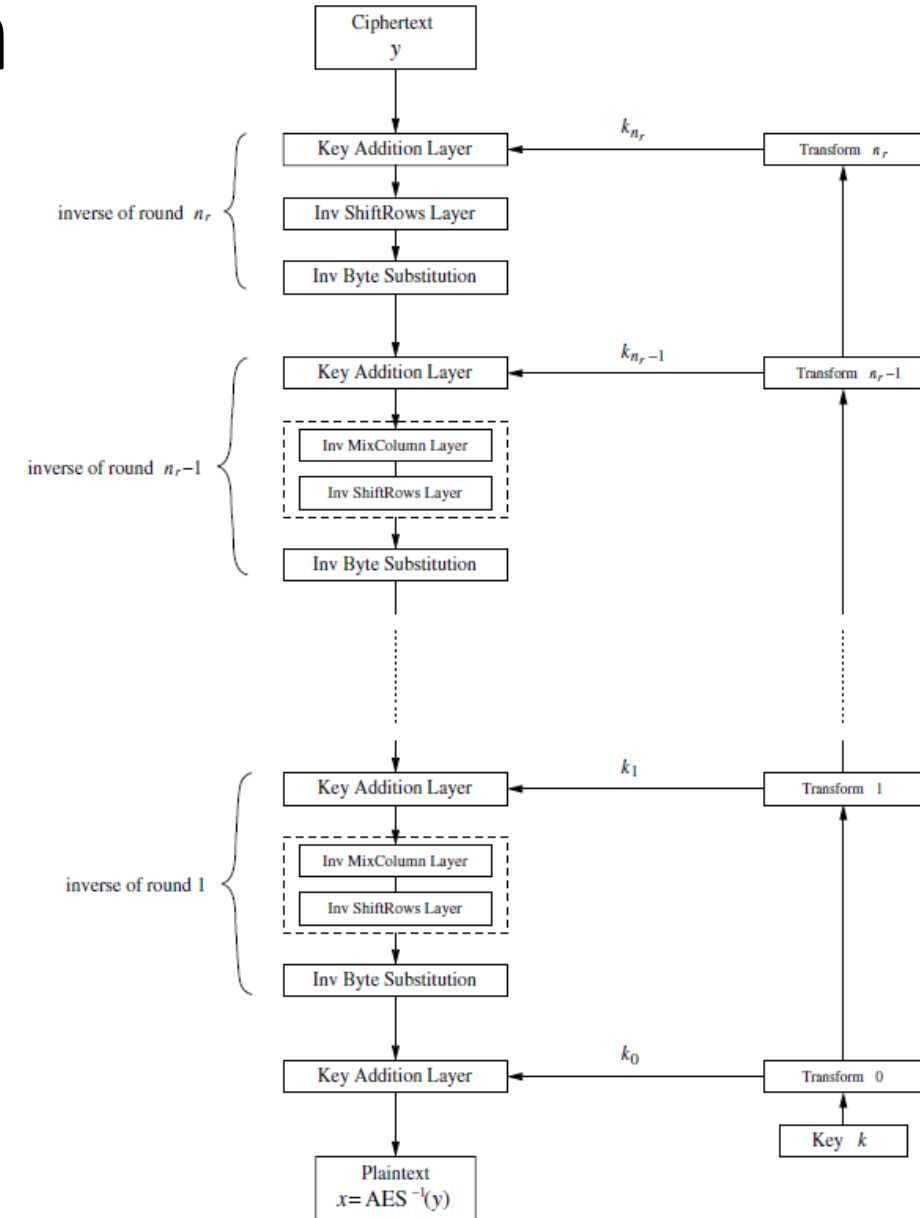
- The two inputs to this layer are the 16-byte state matrix C and the round key k_i (also 16 bytes), where i is the i th round
- The round keys are derived using the key schedule
- The two inputs are combined using a bitwise XOR operation
 - This is equivalent to addition in the Galois extension field $GF(2^8)$
 - Recall that bitwise XOR operation is also equivalent to addition modulo 2, which means that it is also addition in the Galois field $GF(2)$

The AES algorithm: key schedule

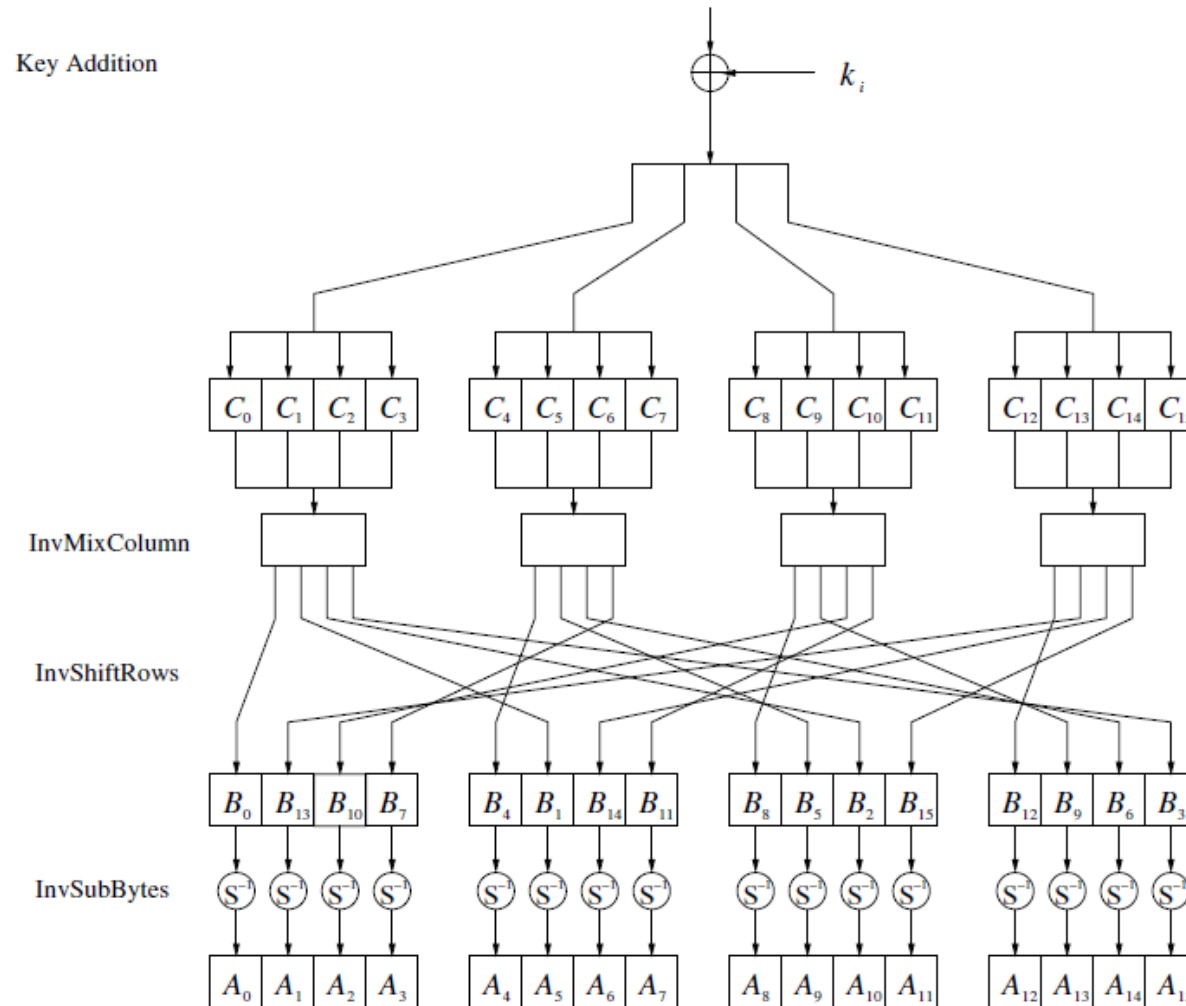
- The key schedule takes the original input key (128-bit, 192-bit or 256-bit) and derives the round keys k_i for $i = 0, 1, \dots, n_r$
 - 128-bit key length — $n_r = 10$ (11 round keys)
 - 192-bit key length — $n_r = 12$ (13 round keys)
 - 256-bit key length — $n_r = 14$ (15 round keys)
- The round keys are computed iteratively
 - i.e. round key k_i must first be computed in order to derive round key k_{i+1} , and so on
- Different key schedules for the three different AES key sizes (128-bit, 192-bit or 256-bit)

The AES algorithm: decryption

- AES is not based on a Feistel network, so all layers need to be inverted for the decryption process (and executed in reverse order)
 - Byte Substitution layer becomes Inv Byte Substitution layer
 - ShiftRows sublayer becomes Inv ShiftRows
 - MixColumn sublayer becomes Inv MixColumn
 - Key Addition layer remains unchanged, because the XOR operation is its own inverse
- A reversed key schedule is also required



The AES algorithm: decryption (single round)



Block cipher modes


- So far, we have discussed the DES and AES block ciphers with respect to their operations on one block of the plaintext
 - What about encrypting a long plaintext consisting of multiple blocks?
- It turns out that there are several ways to encrypt a long plaintext with a block cipher
- One obvious way: Apply the block cipher to each block of the plaintext individually (after padding)
 - This is the **Electronic Code Book** mode (ECB)

Block cipher modes

- Other modes include:
 - Cipher **B**lock **C**haining mode (CBC)
 - Cipher **F**eed**b**ack mode (CFB)
 - **O**utput **F**eed**b**ack mode (OFB)
 - **C**ounter mode (CTR)
- ECB and CBC are actual block ciphers
- CFB, OFB and CTR use a block cipher as a building block for a *stream* cipher
 - These modes use the block cipher to generate a key stream s_i
- Let's discuss each of these five modes in more detail

Block cipher modes: ECB

- ECB mode is the most straightforward way to encrypt a message using a block cipher
 - Requires that the length of the plaintext be an exact multiple of the input block size of the block cipher used (so padding is needed)
- Encryption and decryption in ECB mode:



Let $e()$ be a block cipher in ECB mode with block size b and key k , and x_i and y_i be bitstrings of length b

Encryption: $y_i = e_k(x_i)$, for $i \geq 1$

Decryption: $x_i = e_k^{-1}(y_i) = e_k^{-1}(e_k(x_i))$, for $i \geq 1$

Block cipher modes: comments on ECB

- Advantages of ECB:
 - Can be parallelized, suitable for high-speed implementations
 - Block synchronization between the sending and receiving parties is not necessary
- Disadvantage of ECB:
 - Encrypts in a very deterministic manner; i.e. assuming the key does not change, encrypting a plaintext block always result in the same ciphertext block

Block cipher modes: comments on ECB

- A block cipher in ECB mode encrypts in a very deterministic manner
- As such, it is vulnerable to traffic analysis – an attacker can recognize if the same plaintext message has been sent twice by simply looking at the ciphertext (note: the attacker does **not** know the contents of the plaintext)
- Also, an attacker can *reorder* the ciphertext blocks transmitted by the sending party, which might lead to a valid (but different) plaintext
- So a block cipher in ECB mode is vulnerable to substitution attacks
 - Attacker manipulates the ciphertext by replacing some ciphertext blocks with other ciphertext blocks that are observed and/or manipulated by the attacker
 - Note that this is **not** an attack that breaks the block cipher itself – this is a message *integrity* issue (recall Lectures 3 & 4)

Block cipher modes: comments on ECB

- A block cipher in ECB mode encrypts in a very deterministic manner
 - Identical plaintext values are mapped to the same ciphertext values (on a per-block basis)
- E.g. Encrypting a bitmap image using AES in ECB mode with a 256-bit key:

**CRYPTOGRAPHY
AND
DATA SECURITY**




Block cipher modes: CBC

- CBC mode addresses the security issue of the ECB mode
- Encryption is randomized using an initialization vector (IV)
 - The IV is typically randomly generated and is used only once (ideally), thus it is also known a *nonce* (i.e. “number only used once”)
 - Note that the IV does **not** need to be kept secret; it can be transmitted over an insecure channel between the sending and receiving parties
 - Choose a new IV each time the message is encrypted
- The ciphertext block y_i depends not only on plaintext block x_i , it also depends on all previous plaintext blocks x_{i-1}, \dots, x_1

Block cipher modes: CBC

- Encryption and decryption in CBC mode:



Let $e()$ be a block cipher in CBC mode with block size b and key k , and x_i and y_i be bitstrings of length b , and IV be a nonce of length b

Encryption (first block): $y_1 = e_k(x_1 \oplus IV)$ *reused over.*

Encryption (subsequent blocks): $y_i = e_k(x_i \oplus y_{i-1})$, for $i \geq 2$

Decryption (first block): $x_1 = e_k^{-1}(y_1) \oplus IV$

Decryption (subsequent blocks): $x_i = e_k^{-1}(y_i) \oplus y_{i-1}$, for $i \geq 2$

Block cipher modes: comments on CBC

- Assuming the IV is changed every time encryption is done, then substitution attacks will no longer work
 - Encrypting the same plaintext twice results in different ciphertexts
- However, the attacker can still cause mischief: he can substitute the original ciphertext blocks with garbage blocks (though he can no longer effectively manipulate the original ciphertext blocks)

Block cipher modes: comments on CBC

- Basic point is that encryption itself is insufficient – message integrity also needs to be protected; hence the need for hashes and MACs as mentioned earlier
- On that note, the CBC mode *could* be used to generate a MAC – Alice can transmit the last ciphertext block y_n as a MAC (instead of all of the y_i blocks), along with the plaintext message x
 - This is known as CBC-MAC
 - To verify message integrity, Bob repeats the **encryption** process of the block cipher in CBC mode on x to obtain y_n' , then compares y_n' with the received y_n
 - However, this is not a really secure way to generate a MAC

Block cipher modes: comments on CBC

- Why is the CBC-MAC mode not a secure MAC function?

- Let $x = (x_1, x_2)$ and $h_i = e_k(h_{i-1} \oplus x_i)$, with $h_0 = IV$

So $h_1 = e_k(h_0 \oplus x_1) = e_k(IV \oplus x_1)$, and $h_2 = e_k(h_1 \oplus x_2)$

- Now let's craft an alternate plaintext $x' = (x'_1, x'_2)$ such that:

$$x'_1 = IV \oplus h_1 \oplus x_2 \text{ and } x'_2 = h_2 \oplus h_1 \oplus x_2$$

- Then $h'_1 = e_k(h_0 \oplus x'_1) = e_k(IV \oplus x'_1) = e_k(IV \oplus IV \oplus h_1 \oplus x_2) = e_k(h_1 \oplus x_2)$

$$\text{and } h'_2 = e_k(h'_1 \oplus x'_2) = e_k[e_k(h_1 \oplus x_2) \oplus h_2 \oplus h_1 \oplus x_2]$$

$$= e_k[e_k(h_1 \oplus x_2) \oplus e_k(h_1 \oplus x_2) \oplus h_1 \oplus x_2] = e_k(h_1 \oplus x_2) = h_2$$

- So, we have $x' \neq x$ but $h'_2 = h_2$ and we obtain a collision; this MAC function does not meet the second preimage resistance requirement

CBC-MAC \rightarrow another example of 2nd preimage collision.

let $x = x_1 \rightarrow$ one block message
with $h_i = e_k(h_{i-1} \oplus x_i)$, $h_0 = IV$

$$h_1 = e_k(h_0 \oplus x_1) \\ = e_k(IV \oplus x_1)$$

lets make alternate msg: $x' = (x'_1, x'_2) \rightarrow$ 2 block message.


$$\text{with } x'_1 = x_1 \\ x'_2 = IV \oplus h_1 \oplus x_2$$

$$\Rightarrow h'_1 = e_k(h_0 \oplus x'_1) = e_k(IV \oplus x_1) = h_1 \\ h'_2 = e_k(h'_1 \oplus x'_2) = e_k(h_1 \oplus IV \oplus h_1 \oplus x_2) \\ = e_k(IV \oplus x_2) = h_2$$

we have $h'_2 = h_2$
 \rightarrow MAC are the same for x and x' .

Block cipher modes: comments on CBC

- A better way to use a block cipher as a hash function (Davies-Meyer):




- This function sequentially utilizes the plaintext blocks x_i as the keys k_i (with a b -to- m mapping function, where m is the key length of the block cipher)
- $h_i = h_{i-1} \oplus e_{x_i}(h_{i-1})$, where $e_{x_i}()$ is the encryption function of the block cipher

Block cipher modes: OFB ↗ output feedback mode.

- OFB mode uses a block cipher as a building block for a stream cipher
 - It uses the block cipher to generate a key stream s_i
 - The key stream is **not** generated bitwise; it is generated in a block-wise manner
- An IV (should be a nonce) is encrypted using the block cipher and the cipher output is the first set of b key stream bits
- The cipher output is also fed back into the block cipher to generate the next block of b key stream bits
- This process is repeated as necessary
- The actual encryption of the plaintext is handled by the XOR function, just like in a regular stream cipher – the plaintext bits are XOR-ed with the key stream bits

Block cipher modes: OFB

- Encryption and decryption in OFB mode:



Let $e()$ be a block cipher with block size b and key k , with x_i , y_i and s_i being bitstrings of length b , and let IV be a nonce of length b

Encryption (first block): $s_1 = e_k(IV)$, $y_1 = s_1 \oplus x_1$

Encryption (subsequent blocks): $s_i = e_k(s_{i-1})$, $y_i = s_i \oplus x_i$, for $i \geq 2$

Decryption (first block): $s_1 = e_k(IV)$, $x_1 = s_1 \oplus y_1$

Decryption (subsequent blocks): $s_i = e_k(s_{i-1})$, $x_i = s_i \oplus y_i$, for $i \geq 2$

Block cipher modes: comments on OFB


- Like CBC mode, the encryption in OFB mode is nondeterministic
 - i.e. encrypting the same plaintext twice results in different ciphertexts
- Still need to protect message integrity, as before
- Block cipher computations are independent of the plaintext bits, so several blocks of the key stream bits can be computed in advance

Block cipher modes: CFB

- Like OFB mode, CFB mode also uses a block cipher as a building block for a stream cipher; the key stream is generated block-wise
- An IV is encrypted using the block cipher and the cipher output is the first set of b key stream bits (just like OFB)
- The difference now is that ciphertext output y_1 (which is the output of the **stream cipher**) is fed back into the block cipher to generate the next block of b key stream bits
- The next ciphertext output y_2 is then fed back into the block cipher to generate the following block of b key stream bits
- This process is repeated as necessary

Block cipher modes: CFB

- Encryption and decryption in CFB mode:



Let $e()$ be a block cipher with block size b and key k , with x_i and y_i being bitstrings of length b , and let IV be a nonce of length b

Encryption (first block): $y_1 = e_k(IV) \oplus x_1$

Encryption (subsequent blocks): $y_i = e_k(y_{i-1}) \oplus x_i$, for $i \geq 2$

Decryption (first block): $x_1 = e_k(IV) \oplus y_1$

Decryption (subsequent blocks): $x_i = e_k(y_{i-1}) \oplus y_i$, for $i \geq 2$

Block cipher modes: comments on CFB


- The encryption in CFB mode is nondeterministic, like OFB
 - i.e. encrypting the same plaintext twice results in different ciphertexts
- Still need to protect message integrity, as always

Block cipher modes: CTR

- CTR mode also uses a block cipher as a building block for a stream cipher; the key stream is generated block-wise
- The input to the block cipher is a counter whose value changes every time the block cipher computes a new key stream block
- Care needs to be taken with the initialization of the input to the block cipher, to avoid using the same input value twice
 - This can be achieved by concatenating some *IV* (that is a nonce with a length smaller than the block size of the block cipher) with the counter value
 - Length of the *IV* and counter value is equal to the block size
 - E.g. Assuming AES is used as the block cipher, choose 96 bits for the *IV* and 32 bits for the counter value, for a total of 128 bits
 - Counter value is initialized to zero
 - Each time a block is encrypted, the *IV* stays the same while the counter is incremented

Block cipher modes: CTR

- Encryption in CTR mode:



Let $e()$ be a block cipher with block size b and key k , with x_i and y_i being bitstrings of length b , and let $IV || CTR_i$ denote the concatenation of the initialization value IV with the counter value CTR_i ; this concatenation is a bitstring of length b

Encryption: $y_i = e_k(IV || CTR_i) \oplus x_i$, for $i \geq 1$

Decryption: $x_i = e_k(IV || CTR_i) \oplus y_i$, for $i \geq 1$

Block cipher modes: comments on CTR

- The encryption in CTR mode is nondeterministic, like OFB and CFB
 - i.e. encrypting the same plaintext twice results in different ciphertexts
- Still need to protect message integrity
- CTR mode can be parallelized, unlike OFB and CFB
 - E.g. two block cipher devices running in parallel – the first device encrypts with $IV || CTR_1$ and the second one encrypts with $IV || CTR_2$
 - Then after they are done, the first device encrypts with $IV || CTR_3$ and the second one encrypts with $IV || CTR_4$, and so on
 - Good for applications with high throughput requirements